

OpenMP Parallel Programming Project Report

Rishik Sarkar (rs1894)

Experimental Results Table:

	1	2	3	4	5	6	7	8	9	10	Avg.	Med.	Std.
s-2	145.23	145.19	144.64	145.85	145.71	144.23	144.84	148.31	143.90	143.05	145.10	145.02	1.339
sol	146.18	145.11	143.55	144.77	149.18	146.83	188.94	146.19	142.42	144.60	149.78	145.65	13.172
	1	2	3	4	5	6	7	8	9	10	Avg.	Med.	Std.
f-2	137.84	136.91	137.06	132.02	134.43	135.61	136.47	137.19	136.90	139.05	136.35	136.91	1.853
sol	137.23	137.69	137.42	156.07	153.51	138.24	140.93	137.91	136.06	139.86	141.49	138.08	6.797
	1	2	3	4	5	6	7	8	9	10	Avg.	Med.	Std.
s-4	69.29	68.20	68.99	68.06	67.88	72.47	68.90	68.86	68.39	67.45	68.85	68.63	1.322
sol	68.58	68.73	70.33	67.96	69.53	68.11	67.68	70.50	73.45	67.69	69.26	68.66	1.704
	1	2	3	4	5	6	7	8	9	10	Avg.	Med.	Std.
f-4	69.06	69.59	67.80	67.66	69.40	69.62	69.02	66.66	69.45	68.71	68.70	69.04	0.947
sol	68.96	69.57	68.13	68.62	68.48	69.50	69.28	69.01	72.59	69.13	69.33	69.07	1.169

Notes:

- The **rows** represent particular **code versions** with their respective solution file right below. The file names are as follows:
 - s-2**: *spell_t2_singleloop*
 - f-2**: *spell_t2_fastest*
 - s-4**: *spell_t4_singleloop*
 - f-4**: *spell_t4_fastest*
- The first ten **columns** represent the **execution times (in ms)** of each experiment conducted on the specified code version of the row. The last three columns represent the **average (mean), median, and standard deviation** of the entries of all 10 experiments in each respective row.
- The execution times have been rounded to 2 decimals (except standard deviation, which is rounded to 3 decimals) in order to fit the data within the table
- All the experiments were conducted on **iLab1.cs.rutgers.edu**, and attempted to spell check the word “**principles**”

Parallelization Strategy:

- I chose to parallelize the **outer loop** in the nested for-loop that applies every hash function in hf to every word in the dictionary of words and fills the bit vector (bv). I picked this loop because:
 - Each iteration is independent, so we can distribute them over multiple threads
 - Parallelizing the outer loop is better in this case since it seems to have a lower overhead than the inner loop
- For most of the programs, I used the following statement:

```
#pragma omp parallel for private(i, j, hash) schedule(guided) shared(hf, bv_size, bv)
```

 - Parallelizes the outer for loop
 - Creates private copies of the local variables i, j, and hash in each thread
 - Divides the iterations into similar sized chunks that decrease in size later in order to provide optimal load balancing. I mainly used this because it led to much better performance than static or dynamic scheduling
 - Shares the variables/data structures hf, bv_size, and bv across all threads